

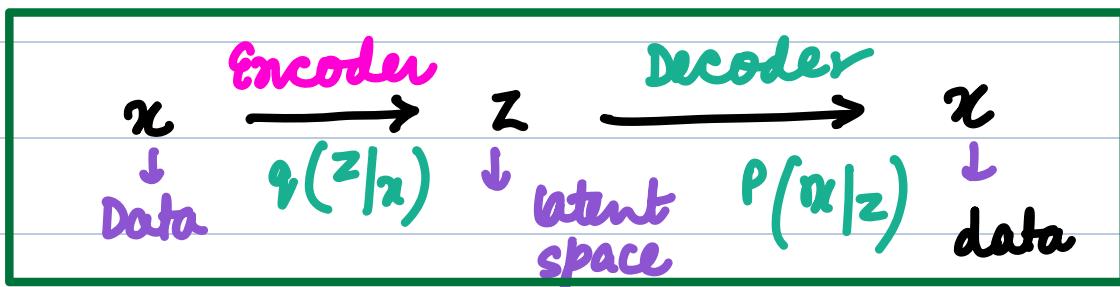
Denoising Diffusion Probabilistic

models
(DDPM)

Given data $D = \{x_0\} \sim p_{x_0}$ (data distribution)

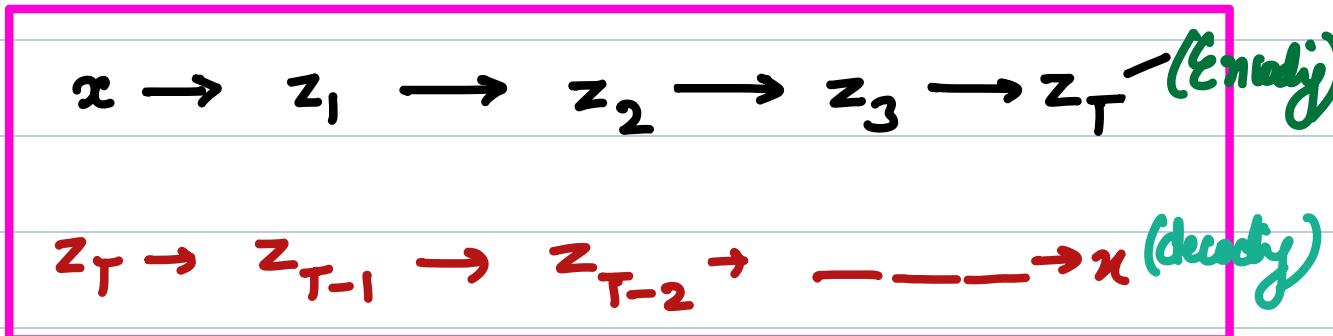
Goal → To learn to sample from
 p_{x_0}

(special case of Hierarchical VAE)



—VAE

HVAE



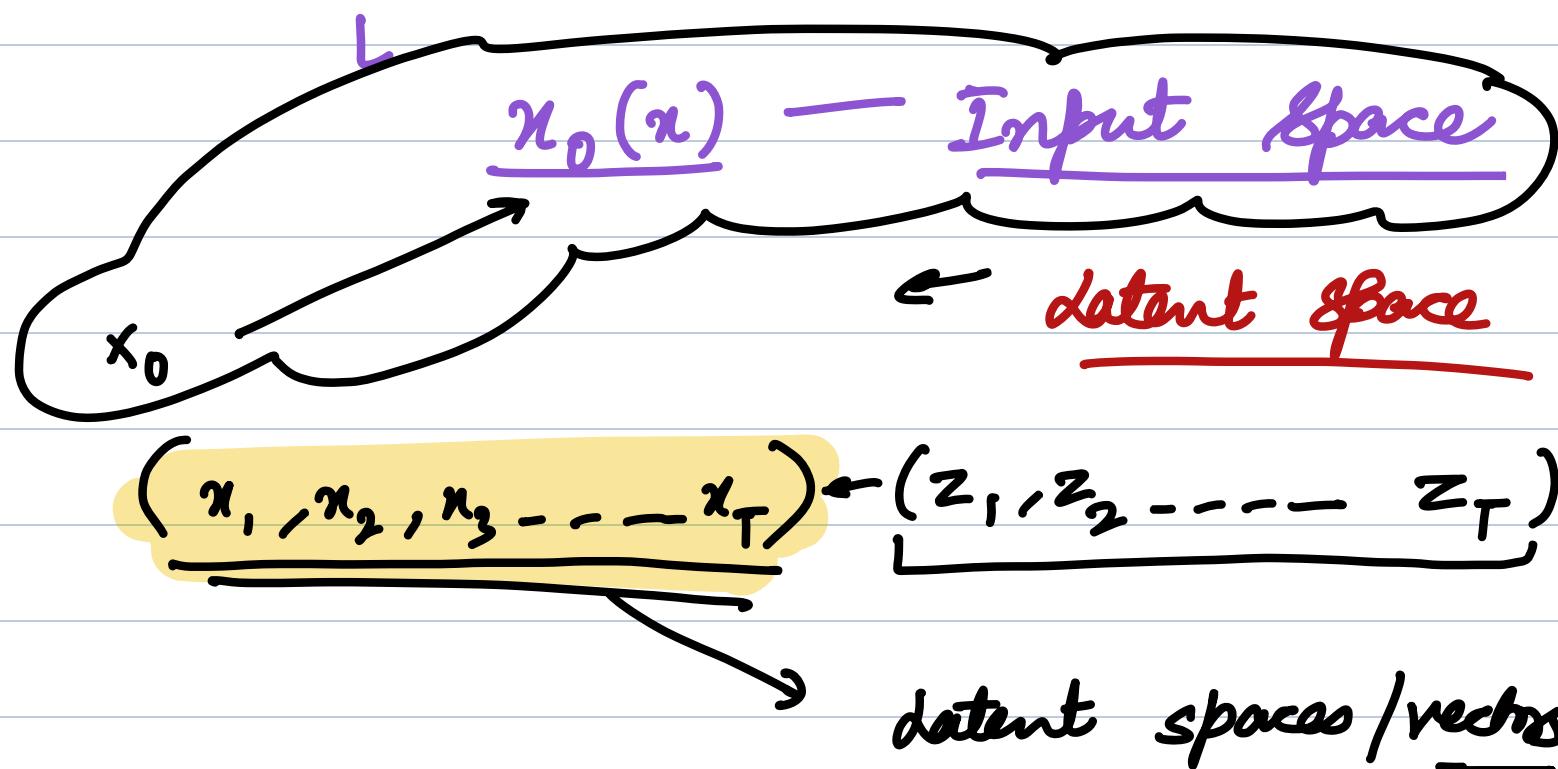
DDPM

HVAE with
1) Multiple
latent
spaces

Dimensionality
of latent
space same as
data space

Encoding
process is
non-learnable
/ fixed

VAE $\rightarrow q_{\phi}(z|x) \rightarrow$ learnt ✓
 Both encode & decoding learnt.
 DDPM $\rightarrow q(z|x)$ — fixed / not learnt
 only decoding is learned



Forward Process

$x_0 \rightarrow x_1 \rightarrow x_2 \rightarrow x_3 \dots x_T$
 input
 latent vectors/
 spaces
 corresponding to x_0

$x_1 = \sqrt{\alpha_1} x_0 + \sqrt{1-\alpha_1} \epsilon_1$
 noise

α = scalar

where $\epsilon_1 \sim N(0, I)$

$$x_2 = \sqrt{\alpha_2} x_1 + \sqrt{1-\alpha_2} \epsilon_2$$

where $\epsilon_2 \sim N(0, I)$

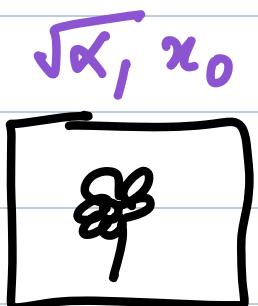
$$x_3 = \sqrt{\alpha_3} x_2 + \sqrt{1-\alpha_3} \epsilon_3$$

$$x_t = \sqrt{\alpha_t} x_{t-1} + \sqrt{1-\alpha_t} \epsilon_t$$

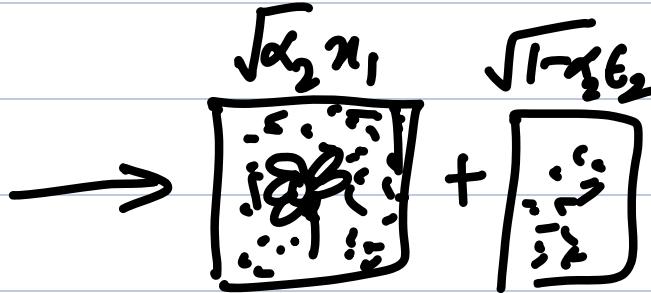
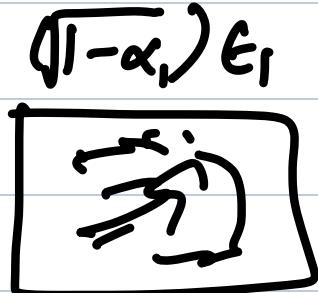
$\alpha_1, \alpha_2, \alpha_3, \dots$

$\alpha_t \sim$ are
fixed scalars

$[0, 1]$



+



— — — — —



Forward Process

$$q(x_t | x_{t-1}) = N(x_t; \sqrt{\alpha_t} x_{t-1}, (1-\alpha_t)I)$$

(Conditional forward encoding)

Output

Mean

Variance

β_t : small variance schedule

$$\alpha_t + \beta_t = 1$$

$$\boxed{\beta_t = 1 - \alpha_t}$$

Define the Model —

$$P_\theta(x_0, x_1, x_2, \dots, x_T) = P_\theta(x_T) \prod_{t=1}^T P_\theta(x_t | x_{t-1})$$

data *latent variable*

Reverse Process

$$P_\theta(x_T) \prod_{t=1}^T P_\theta(x_{t-1} | x_t)$$

Mean

where

$$p_{\theta}(x_{t-1}, x_t) \stackrel{?}{=} N(x_{t-1}, h_{\theta}(x_t), \Sigma_{\theta}(z))$$

output *variance*

ELBO optimization
↓ VAE

VAE

$$\log p_{\theta}(x) = \log \int_z p_{\theta}(x, z) dz$$

$$\text{loss} \approx J_{\theta}(q_{\phi}) = E_{q_{\phi}(z|x)} \frac{\log p_{\theta}(x, z)}{q_{\theta}(z|x)}$$

evidence
lower bound
(ELBO)

In DDPM

$$\approx J_{\theta}(q_r) \stackrel{\text{DDPM}}{=} E_{q_r(x_1, x_2, \dots, x_T | x_0)} \log \frac{p_{\theta}(x_0, x_1, \dots, x_T)}{q_r(x_1, x_2, \dots, x_T | x_0)}$$

$$x_{1:T} = (x_1, x_2, \dots, x_T)$$

$$x_{0:T} = (x_0, x_1, x_2, \dots, x_T)$$



$$J_\theta(q)^{\text{DDPM}} = E_q[x_{1:T} | x_0] \log \frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)}$$

[Optimizing the ELBO for DDPM]

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p_\theta(x_{t-1} | x_t)$$

$$p(x_T) = \mathcal{N}(0, I)$$

Similarly,

$$q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$$

$$q(x_{1:T} | q_0) = q\left(\frac{x_1}{x_0}\right) q(x_2 | x_0)$$

$$q(x_3 | x_2, x_0)$$

$$= q(x_T | x_{T-1}, x_{T-2}, \dots, x_0)$$

$$= \prod_{t=1}^T q(x_t | x_{t-1})$$

encoding
Forward
encoder