

CS5202 Practice Questions (Set-I)

Spring 2026

1. Generative Models comparison

- (a) Suppose we care about having a model with a tractable inverse mapping from data space to latent space. Which models provide this property by design?
- (b) Suppose we want to minimize a divergence that is not likelihood-based, but instead depends on distinguishing real and fake samples. Which class of models does this correspond to?
- (c) Suppose we want to compute the exact log-likelihood gradient with respect to model parameters for a given datapoint without using variational approximations. Which of these models allow this? Briefly explain.
- a) Autoregressive models b) Latent variable models c) Flow-based models d) all of the above
- (d) VAEs optimize:

$$\log p_\theta(x) = \text{ELBO} + D_{KL}(q_\phi(z | x) \| p_\theta(z | x)).$$

If the encoder has insufficient capacity, what effect does this have on generative quality?

2. Latent Variable Model and ELBO

Consider the joint distribution of a latent variable model denoted by $p(x, z)$. The model is capable of sampling only two images $\{x^{(1)}, x^{(2)}\}$, where $x^{(i)} \in \{0, 1\}^{784}$. The latent variable is scalar $z \in \mathbb{R}$. The model is defined as:

$$p(z) = \mathcal{N}(z; 0, 1)$$

$$p(x | z) = \begin{cases} 1 & \text{if } z \geq 0 \wedge x = x^{(1)} \\ 0 & \text{if } z \geq 0 \wedge x \neq x^{(1)} \\ 1 & \text{if } z < 0 \wedge x = x^{(2)} \\ 0 & \text{if } z < 0 \wedge x \neq x^{(2)} \end{cases}$$

- (a) Compute the log-likelihood:

$$\text{like}(x) := \log p(x).$$

What is $\text{like}(x^{(1)})$?

- (b) Compute the posterior distribution $p(z | x^{(1)})$ explicitly.

(c) Let the variational distribution be:

$$q(z) = \mathcal{N}(z; 0, 1).$$

Compute the KL divergence:

$$D_{KL}(q(z) \| p(z | x^{(1)})).$$

(d) Now consider instead:

$$q(z) = \mathcal{N}(z; 1, 1).$$

Compute numerically:

$$P_q(z \geq 0).$$

(e) Using your answer above, compute the ELBO:

$$\text{ELBO}(x^{(1)}; q) = \log p(x^{(1)}) - D_{KL}(q(z) \| p(z | x^{(1)})).$$

Is it finite?

3. Theory Questions

- (a) What is the goal of Generative AI. How does it differ from discriminative modeling?
- (b) Give two historical milestones in the development of modern Generative AI and explain why they were significant.
- (c) Explain why latent variable models are useful for high-dimensional data such as images or speech.
- (d) Suppose that you are modeling medical images where the interpretability of latent factors is important. What advantages do latent variable models provide compared to autoregressive models?
- (e) Why is variational inference often preferred over exact Bayesian inference in deep generative models?
- (f) What are the main limitations of variational inference? Discuss the role of the variational gap.
- (g) Explain the key differences between a standard autoencoder and a VAE.
- (h) In a real-world application such as anomaly detection in industrial equipment, which model would you choose and why?
- (i) Discuss two advantages and two disadvantages of GANs compared to likelihood-based models.
- (j) You are building a system to generate realistic sports player avatars for a video game. Would you choose a GAN, VAE, or diffusion model? Justify your answer.
- (k) What are the main advantages of normalizing flows over VAEs?
- (l) Why are normalizing flows typically harder to scale to very high-resolution images?

- (m) Explain the main idea behind diffusion models.
- (n) Why do diffusion models often produce higher-quality samples than GANs?
- (o) What is the primary computational disadvantage of diffusion models compared to GANs?
- (p) Explain the forward and reverse processes in Denoising Diffusion Probabilistic Models (DDPM).
- (q) In a real-time content generation application (e.g., live video enhancement), what challenges would arise when deploying DDPM models?
- (r) You are designing a generative model for the following tasks:
 - (a) Generating realistic synthetic MRI scans where likelihood evaluation is required for anomaly scoring.
 - (b) Generating creative artwork for marketing campaigns where visual quality is the primary objective. For each case, select an appropriate model class and justify your choice.

4. Numericals

1. Consider a DDPM defined with a forward diffusion process:

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

where $\{\beta_t\}_{t=1}^T$ is a fixed variance schedule.

- (a) Show that the forward process admits the closed-form:

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$$

where

$$\alpha_t = 1 - \beta_t, \quad \bar{\alpha}_t = \prod_{s=1}^t \alpha_s.$$

- (b) Write the full training objective used in DDPM when the model predicts the noise $\epsilon_\theta(x_t, t)$.

Consider a 1-dimensional diffusion process with:

$$\beta_1 = 0.1, \quad \beta_2 = 0.2.$$

Let $x_0 = 2$.

- (a) Compute α_1, α_2 and $\bar{\alpha}_2$.
- (b) Compute the mean and variance of $q(x_2 | x_0)$.

(c) Explain why the reverse distribution

$$q(x_{t-1} \mid x_t)$$

is Gaussian.

(d) If the model perfectly predicts the true noise, what distribution does the generated sample converge to as $T \rightarrow \infty$?