

Generative AI and LLM

Introduction to Generative AI: From Imagination to Impact
CS5202

Course Instructor : Dr. Nidhi Goyal

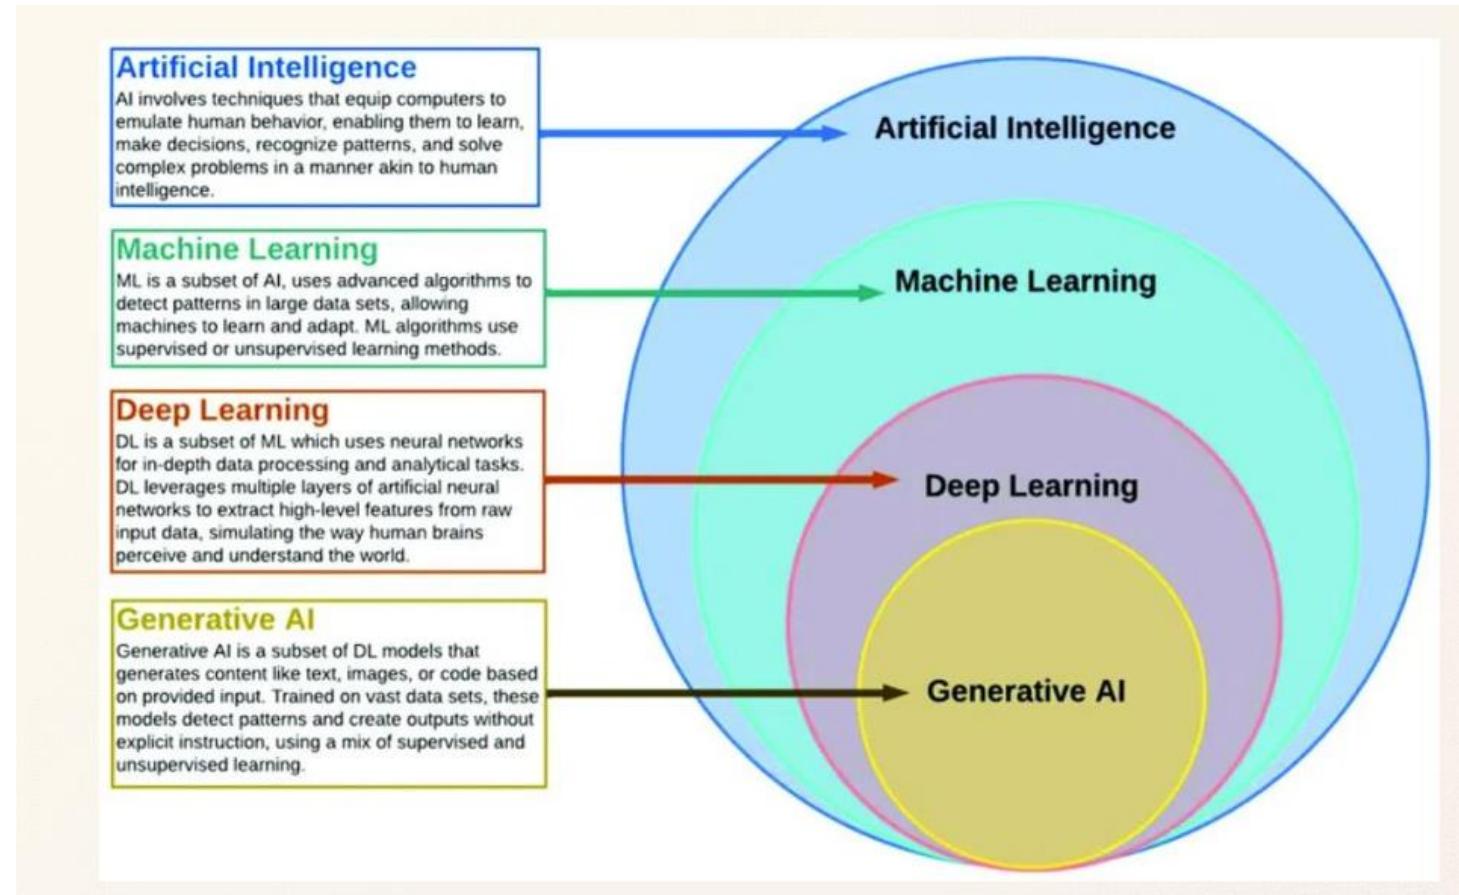
22/1/2026



Lecture Plan

- Impact of Generative AI
- Traditional Ai vs Generative AI
- Evolution of GenAI
- GenAI Taxonomy
- References

The AI Landscape



What If Machines Could Create?

- Write poems, generate code, design drugs
- Diagnose diseases
- Simulate the brain
- Imagine solutions humans haven't thought of

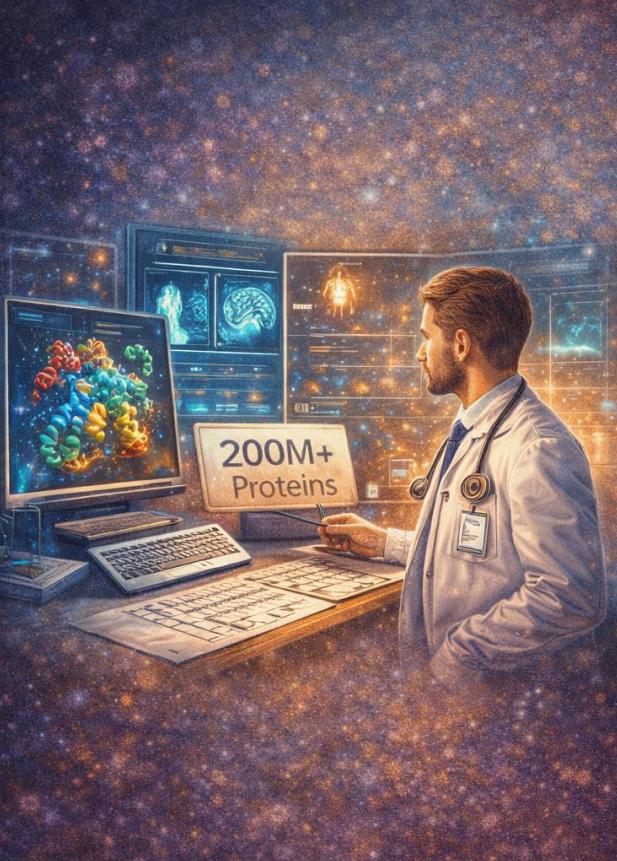
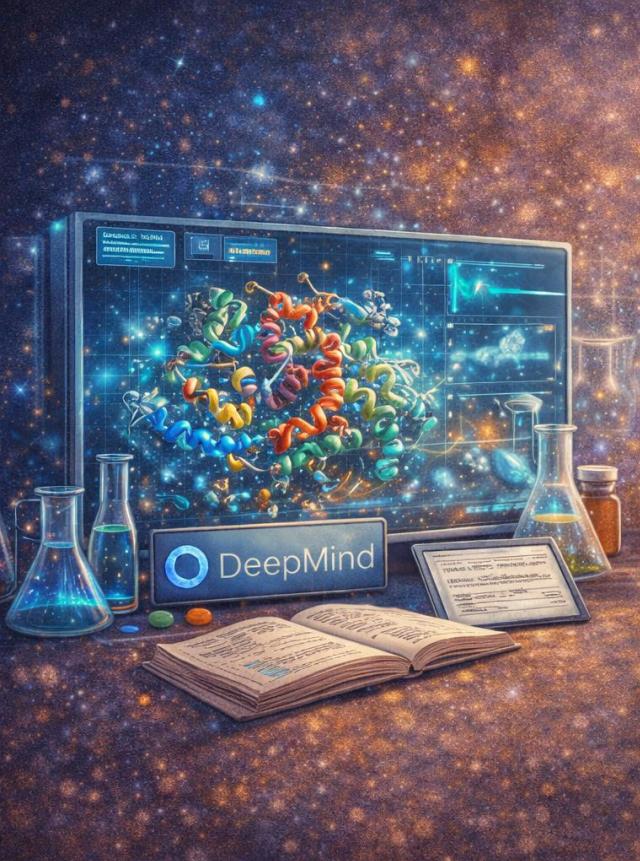
*Generative AI is not about automation — it's about **augmentation of human intelligence***

The GenAI Revolution – ChatGPT

- Single powerful statistic: "**From 0 to 100M users in 2 months - ChatGPT's unprecedented adoption**"
- Brief statement:

Generative AI is not hype—it's already solving problems that were unsolvable 3 years ago





Revolutionizing Medicine and Drug Discovery

- **Key Problems:**
- Slow drug discovery (10–15 years)
- Limited medical expertise in rural areas
- Personalized treatment gaps

GenAI Solutions:

- **DeepMind's AlphaFold**
 - Predicted structures of **200M+ proteins**
 - Accelerated biology & medicine research
- **OpenAI's medical reasoning models**
 - Assist doctors in clinical reasoning
 - Summarize patient histories



22/1/2026



Environment & Climate

Problems:

- Climate modeling is computationally expensive
- Deforestation monitoring is slow
- Energy optimization is complex

GenAI in Action:

- **Google DeepMind AI** for weather & flood prediction
- Generative models for:
 - Climate simulations
 - Smart energy grid optimization
 - Satellite image generation & restoration

Neuroscience & Human Mind

22/1/2026

Challenges:

- Brain is complex, noisy, high-dimensional
- Limited labeled neural data
- Understanding cognition & consciousness

Brain-to-Text

bird



bird

GenAI Applications:

- Generative models for **EEG / fMRI data synthesis**
- Brain-to-text and brain-to-image decoding
- Simulation of neural activity patterns

Brain-to-Image



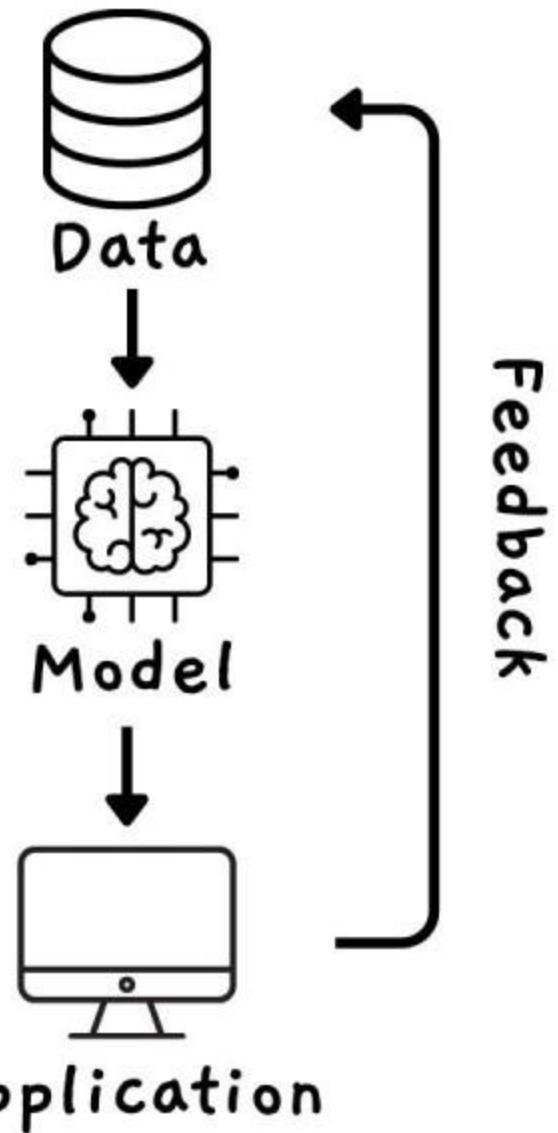
Examples:

- Reconstructing images from brain signals
- Generating synthetic neural data for research

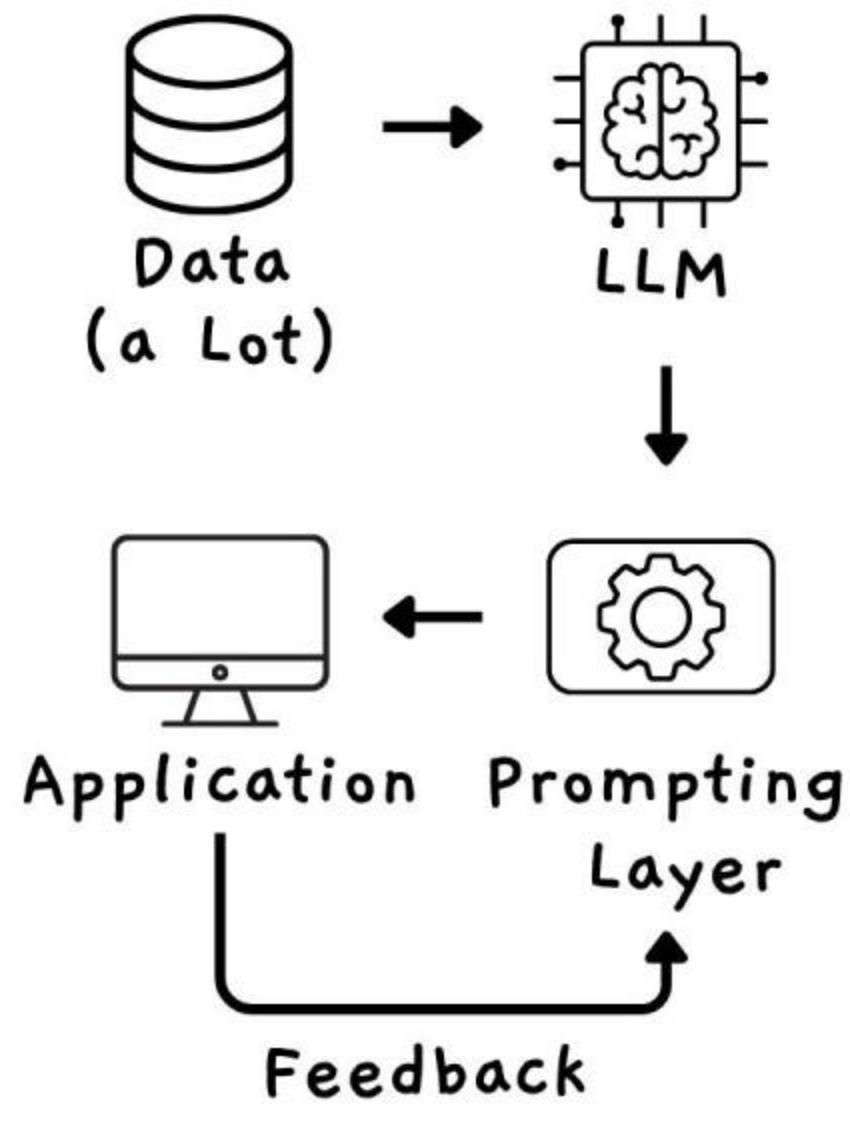
From Impact to Understanding

Now that we've seen what GenAI can achieve, let's explore the fundamentals behind these remarkable capabilities

Traditional (Old-School) AI



Generative AI



What Is Generative AI?

- Generative AI refers to machine learning models that **create new content**-text, images, audio, video, code or data by learning patterns from existing examples.
- **Key distinction:** Unlike traditional AI that classifies or predicts, GenAI *generates* novel outputs that didn't exist before



Discriminative VS Generative AI

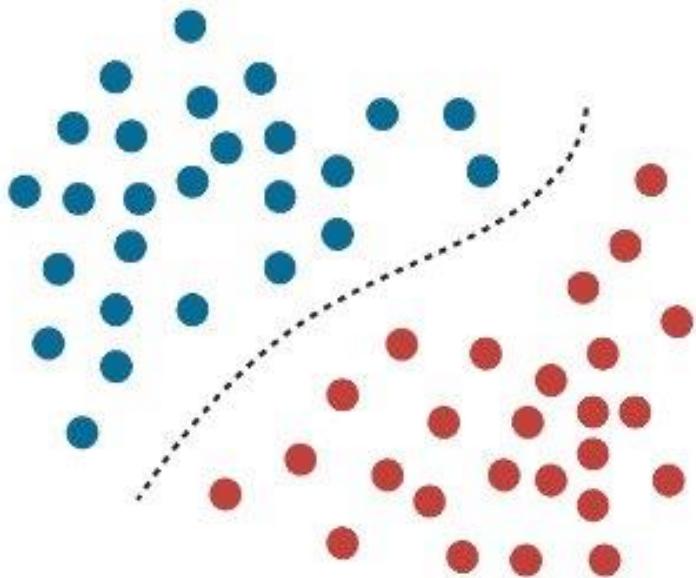
- **Discriminative Models**
- **Goal:** Learn boundaries between classes
- **Question:** "Is this a cat or dog?"
- **Examples:** SVM, Logistic Regression, CNNs for classification
- **Generative Models**
- **Goal:** Learn the data distribution itself
- **Question:** "Can you create a new cat image?"
- **Examples:** GANs, VAEs, Diffusion Models, Transformers

Discriminative vs Generative AI Models

Discriminative (classic)

Predict a label/class given the features of input data

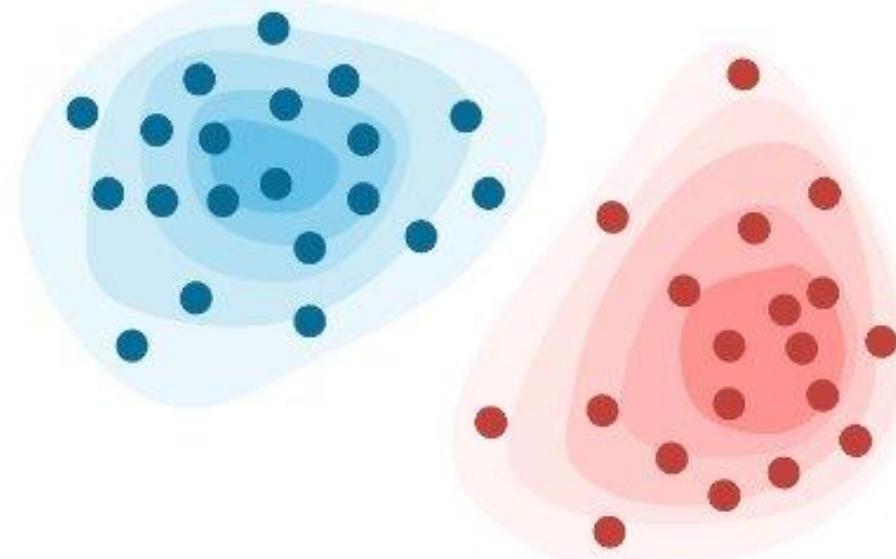
What's learned?: Decision boundary



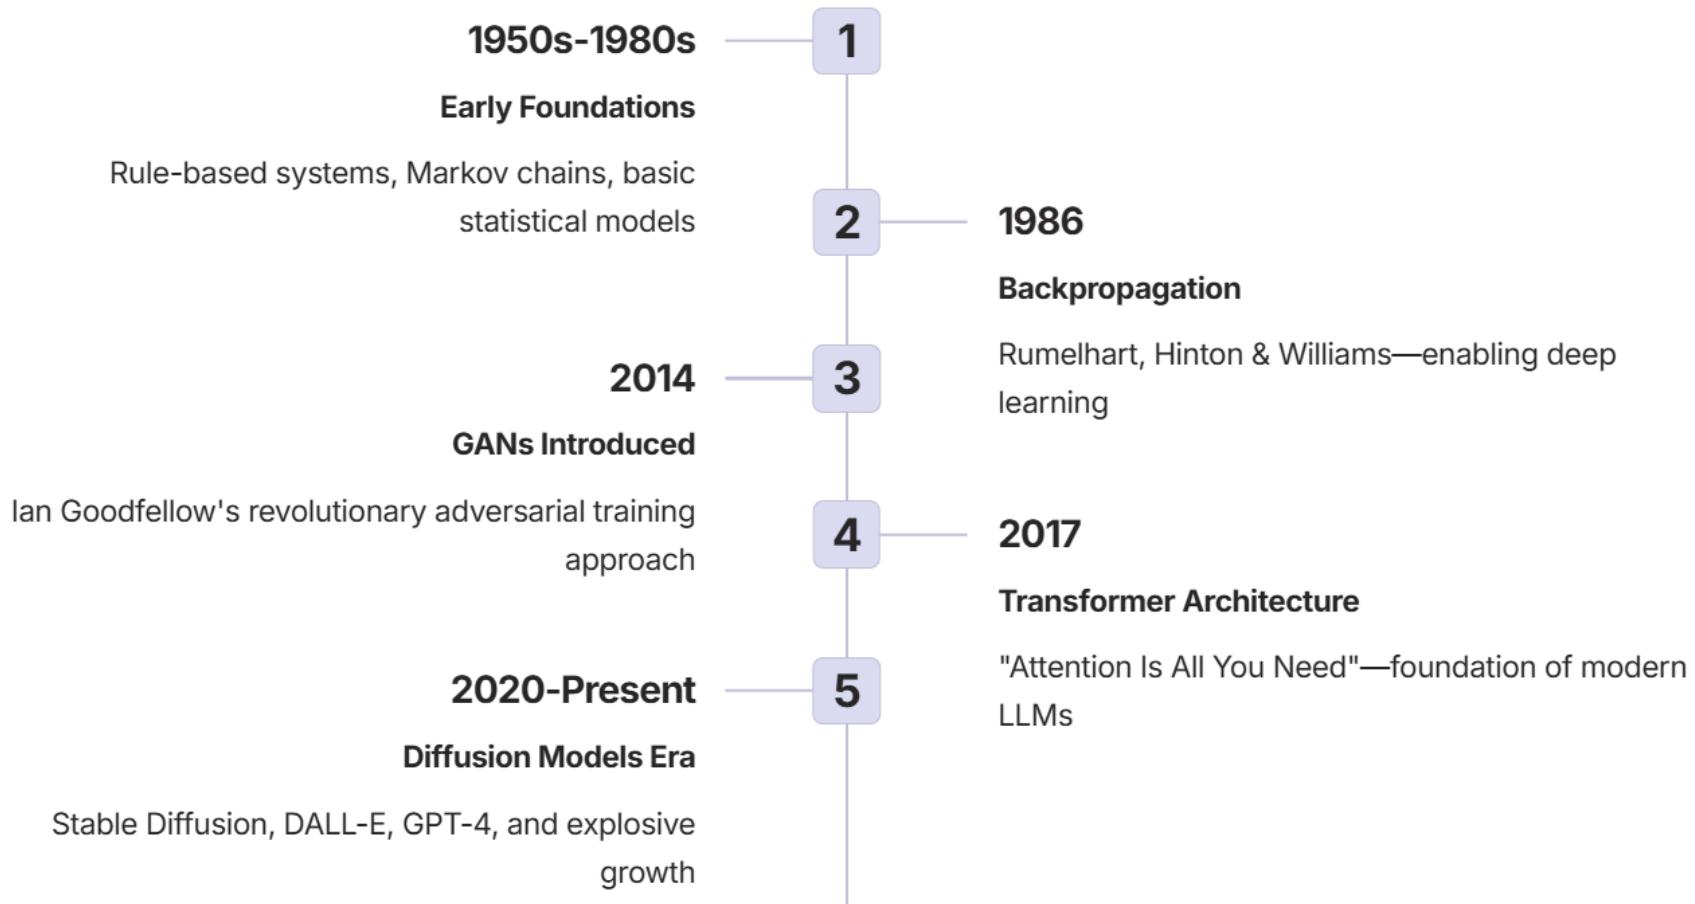
Generative

Abstract underlying patterns in input data in order to generate new content

What's learned?: Probability distributions of the data



Evolution of Generative Models



Key Milestones in the Evolution



Variational Autoencoders (2013)

Kingma & Welling introduce probabilistic latent space learning



GANs Revolution (2014)

Generator vs. Discriminator paradigm produces photorealistic images



Attention Mechanism (2017)

Transformers unlock unprecedented language understanding and generation



Diffusion Models (2020)

Denoising approach achieves state-of-the-art image synthesis



Foundation Models (2022+)

GPT-4, Claude, Gemini—massive scale, multi-modal capabilities

Why Now? Why Not 10 Years Ago?

- Big data 
- GPUs & TPUs 
- Transformers 
- Self-supervised learning
- Open research & scaling laws



Intuition Behind Generative Learning

Core Concept:

- Model sees many examples
- Learns *patterns, structure, uncertainty*
- Can sample new data from learned space

Analogy:

- Like a musician who listens to thousands of songs and composes a new one



Early Era of GenAI (Pre-2010)

Rule-based & probabilistic models

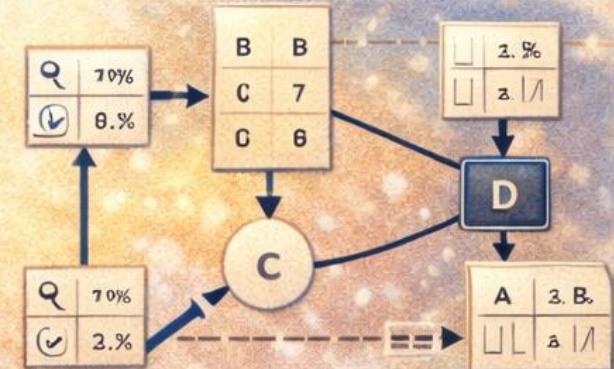
- N-grams
- Hidden Markov Models
- Bayesian networks

Limitations:

- Manual feature design
- Poor scalability
- Weak creativity

Limitations:

- Manual feature design
- Poor scalability
- Weak creativity



Bayesian networks



Rise of Deep Generative Models (2013–2016)

Breakthroughs:

- Autoencoders
- **GANs (2014)**
- Representation learning

Impact:

- Realistic image generation begins
- Data-driven creativity emerges

GANs (2014)

- Generator



Generator



Discriminator



Discriminator

Impact:

- Realistic image generation begins
- Data-driven creativity emerges



GANs & VAEs

GANs:

- Generator vs Discriminator
- Sharp images
- Hard to train

VAEs:

- Probabilistic
- Structured latent space
- Slightly blurry but stable



Transformers Change Everything (2017)

Key Paper: Attention Is All You Need

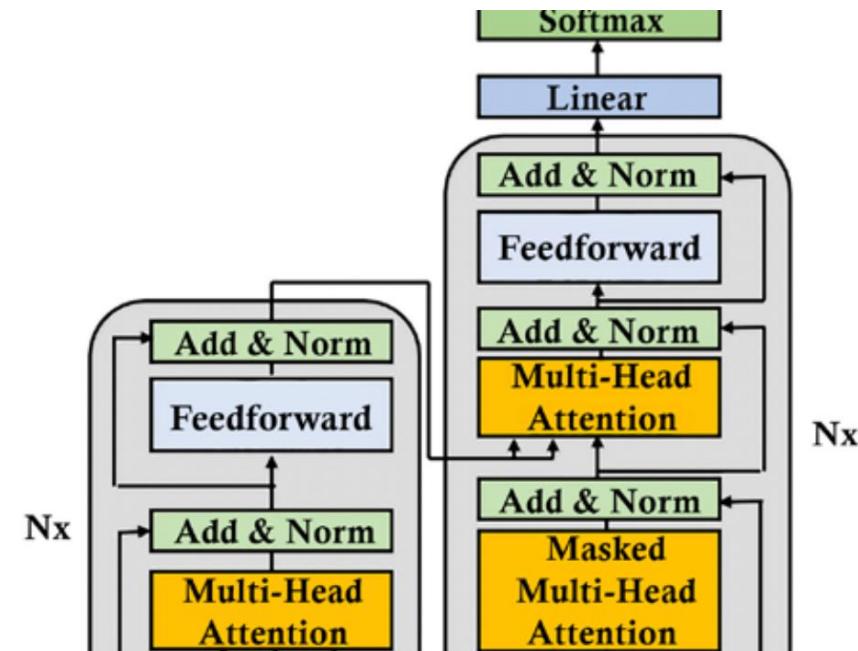


Why Transformers Matter:

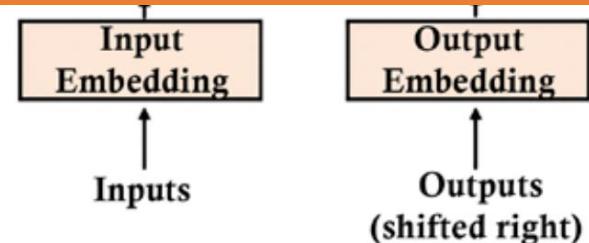
- Parallelism
- Long-range dependencies
- Scaling works

Result:

Foundation models
become possible



Rise of Transformers (2017)



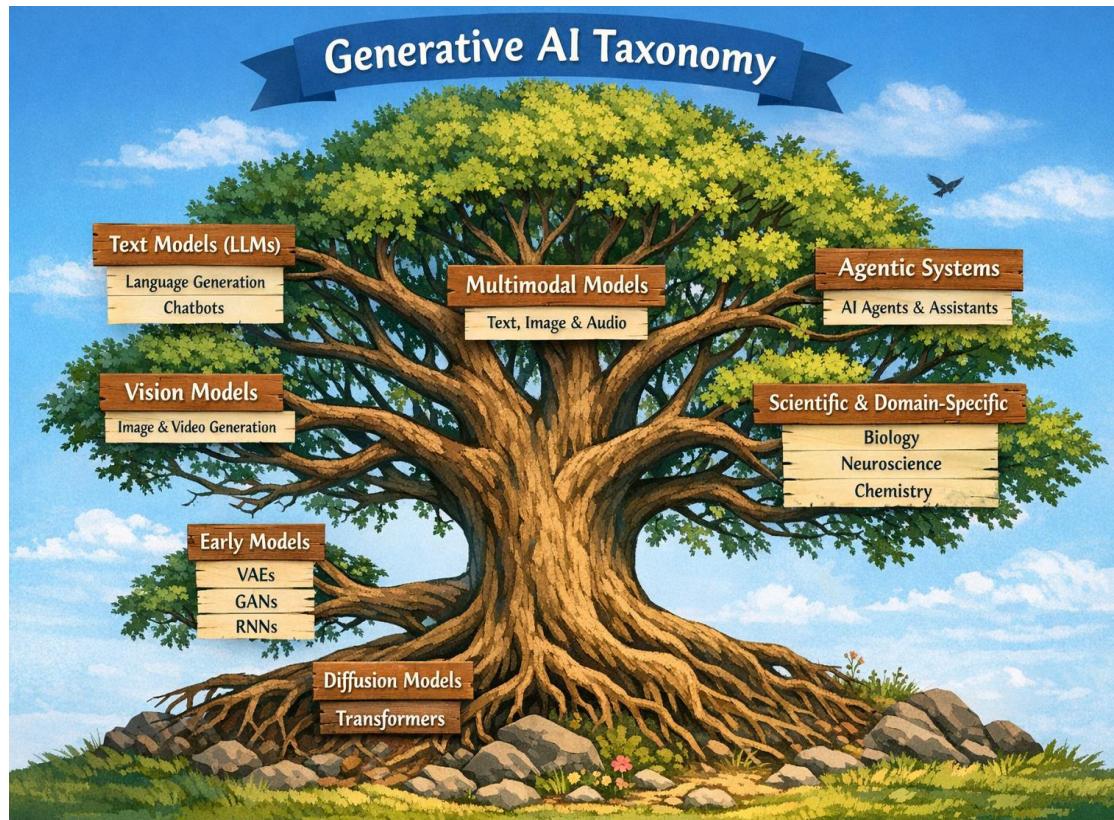
Current GenAI (2020–Now)

Modern Classification:

- Text-only (LLMs)
- Vision-only
- Multimodal (text + image + audio)
- Agentic systems
- Domain-specific GenAI (biology, neuroscience)



GenAI Taxonomy



References

- Vaswani et al., 2017 — *Attention Is All You Need*
<https://arxiv.org/abs/1706.03762>
- Goodfellow et al., 2014 — *Generative Adversarial Networks*
<https://arxiv.org/abs/1406.2661>
- Kingma & Welling, 2013 — *Auto-Encoding Variational Bayes*
<https://arxiv.org/abs/1312.6114>
- Jumper et al., 2021 — *Highly Accurate Protein Structure Prediction with AlphaFold*
<https://www.nature.com/articles/s41586-021-03819-2>
- Shen et al., 2019 — *Deep Image Reconstruction from Human Brain Activity*
<https://www.nature.com/articles/s41593-019-0389-0>